

Smoke Detector! - Binary Classification Using Random Forest and Logistic Regression

Pavlo Mysak

2023-12-13

Objective:

Develop a robust predictive classification model utilizing bio-signals to determine an individual's smoking status from a comprehensive dataset, including age, anthropometric measurements, blood pressure, cholesterol levels, and various biochemical markers. This project employs machine learning techniques (namely, Random Forest and Logistic Regression) to accurately identify smokers and non-smokers, shedding light on the intricate relationship between diverse bio-signals and smoking habits.

Loading the Data, Creating Training & Validation Sets and Quick Summaries

According to the summary of the training data, there are no missing values. We have 22 independent variables available at our disposal.

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(ROSE)
```

```
## Loaded ROSE 0.0-4
```

```
dt <- read.csv('/Users/pavlomysak/Downloads/Kaggle-SmokerBinaryClassification/train.csv')

train.index <- sample(1:nrow(dt), 70000)
train <- dt[train.index,]
valid <- dt[-train.index,]

train$smoking <- factor(train$smoking)
valid$smoking <- factor(valid$smoking)
# remove ID
train <- train[,-1]
valid <- valid[,-1]

summary(train)
```

```

##      age      height.cm.      weight.kg.      waist.cm.
## Min.    :20.0    Min.    :135.0    Min.    : 30.00    Min.    : 51.00
## 1st Qu.:40.0    1st Qu.:160.0    1st Qu.: 60.00    1st Qu.: 77.00
## Median :40.0    Median :165.0    Median : 65.00    Median : 83.00
## Mean    :44.3    Mean    :165.3    Mean     : 67.18    Mean     : 83.02
## 3rd Qu.:55.0    3rd Qu.:170.0    3rd Qu.: 75.00    3rd Qu.: 89.00
## Max.    :85.0    Max.    :190.0    Max.     :130.00    Max.     :127.00
## eyesight.left. eyesight.right. hearing.left. hearing.right.
## Min.    :0.100    Min.    :0.100    Min.    :1.000    Min.    :1.000
## 1st Qu.:0.800    1st Qu.:0.800    1st Qu.:1.000    1st Qu.:1.000
## Median :1.000    Median :1.000    Median :1.000    Median :1.000
## Mean    :1.008    Mean    :1.002    Mean    :1.024    Mean    :1.023
## 3rd Qu.:1.200    3rd Qu.:1.200    3rd Qu.:1.000    3rd Qu.:1.000
## Max.    :9.900    Max.    :9.900    Max.     :2.000    Max.     :2.000
##      systolic      relaxation      fasting.blood.sugar      Cholesterol
## Min.    : 81.0    Min.    : 46.0    Min.    : 46.00    Min.    : 77.0
## 1st Qu.:114.0    1st Qu.: 70.0    1st Qu.: 90.00    1st Qu.:175.0
## Median :121.0    Median : 78.0    Median : 96.00    Median :196.0
## Mean    :122.5    Mean    : 76.9    Mean    : 98.38    Mean    :195.7
## 3rd Qu.:130.0    3rd Qu.: 82.0    3rd Qu.:103.00    3rd Qu.:217.0
## Max.    :213.0    Max.    :121.0    Max.     :375.00    Max.     :393.0
## triglyceride      HDL      LDL      hemoglobin
## Min.    : 8.0    Min.    : 9.00    Min.    : 1.0    Min.    : 5.80
## 1st Qu.: 77.0    1st Qu.: 45.00    1st Qu.: 95.0    1st Qu.:13.80
## Median :115.0    Median : 54.00    Median :114.0    Median :15.00
## Mean    :127.6    Mean    : 55.83    Mean    :114.5    Mean    :14.81
## 3rd Qu.:165.0    3rd Qu.: 64.00    3rd Qu.:133.0    3rd Qu.:15.80
## Max.    :766.0    Max.    :135.00    Max.     :1860.0    Max.     :21.00
## Urine.protein serum.creatinine      AST      ALT
## Min.    :1.000    Min.    :0.1000    Min.    : 6.00    Min.    : 1.00
## 1st Qu.:1.000    1st Qu.:0.8000    1st Qu.: 20.00    1st Qu.: 16.00
## Median :1.000    Median :0.9000    Median : 24.00    Median : 22.00
## Mean    :1.075    Mean    :0.8932    Mean    : 25.51    Mean    : 26.53
## 3rd Qu.:1.000    3rd Qu.:1.0000    3rd Qu.: 29.00    3rd Qu.: 32.00
## Max.    :6.000    Max.    :7.4000    Max.     :656.00    Max.     :2914.00
##      Gtp      dental.caries      smoking
## Min.    : 2.00    Min.    :0.0000    0:39317
## 1st Qu.:18.00    1st Qu.:0.0000    1:30683
## Median :27.00    Median :0.0000
## Mean    :36.27    Mean    :0.1962
## 3rd Qu.:44.00    3rd Qu.:0.0000
## Max.    :766.00    Max.    :1.0000

```

A few of our independent variables are exhibiting suspicious values. triglyceride, LDL, AST, ALT and Gtp have suspiciously high values.

More with Outliers

Generally, we have to be careful with outlier treatment. We don't want to simply omit outliers if they are legitimate observations in our data. We also do not want to blindly alter them (Ex. Winsorization) without domain knowledge. In regard to some machine learning models like Random Forest, we actually want to leave the outliers untouched and perform minimal data cleaning/pre-processing.

To get a better idea of how many outliers we're dealing with, let's iterate through each variable and find how many observations are above the 99th percentile. When we do this, we find that our count of observations over the 99th percentile ranges from around 300 to 700. In a dataset with 70000 observations, this accounts for, at most, less than 1% of the total observations. For this reason, we will disregard these outliers.


```
for( vari in colnames(train)){  
  if(!is.factor(train[[vari]])){  
    n_out = sum(iffelse(train[[vari]] > quantile(train[[vari]], .99),1,0))  
  }  
  cat(vari, ' has', n_out, ' outliers over the 99th percentile.', '\n')  
}
```

```
## age has 281 outliers over the 99th percentile.  
## height.cm. has 18 outliers over the 99th percentile.  
## weight.kg. has 369 outliers over the 99th percentile.  
## waist.cm. has 690 outliers over the 99th percentile.  
## eyesight.left. has 330 outliers over the 99th percentile.  
## eyesight.right. has 285 outliers over the 99th percentile.  
## hearing.left. has 0 outliers over the 99th percentile.  
## hearing.right. has 0 outliers over the 99th percentile.  
## systolic has 666 outliers over the 99th percentile.  
## relaxation has 369 outliers over the 99th percentile.  
## fasting.blood.sugar has 696 outliers over the 99th percentile.  
## Cholesterol has 660 outliers over the 99th percentile.  
## triglyceride has 697 outliers over the 99th percentile.  
## HDL has 592 outliers over the 99th percentile.  
## LDL has 655 outliers over the 99th percentile.  
## hemoglobin has 578 outliers over the 99th percentile.  
## Urine.protein has 249 outliers over the 99th percentile.  
## serum.creatinine has 361 outliers over the 99th percentile.  
## AST has 682 outliers over the 99th percentile.  
## ALT has 679 outliers over the 99th percentile.  
## Gtp has 684 outliers over the 99th percentile.  
## dental.caries has 0 outliers over the 99th percentile.  
## smoking has 0 outliers over the 99th percentile.
```

Distribution of our Dependant Variable

We see that our data has a good number of both smokers and non-smokers. This will give our models sufficient information on both categories to train on.

```
plot(train$smoking)
```




Random Forest

To begin our modelling process, let's initialize a basic random forest model to get an idea of how this algorithm will react to our data and which variables are most important. Random Forest is a tree-based machine learning algorithm that grows a forest of (weak-ish) classification trees using a process called 'bagging'. Bagging, or Bootstrap Aggregating is an 'ensemble learning method' that selects n uniform random samples with replacement from a training set and fits them to n classification models, respectively. In our case, it would create n 'weak' random forest trees. These trees are considered 'weak' because they may have limited predictive power on their own. The strength of the Random Forest algorithm materializes when we combine the classifications of multiple trees. To return our ultimate classification (Ensemble Classifier), each tree 'votes' on a classification. The 'forest' returns the classification that has received the most votes from the population of trees.

```
rf_init <- randomForest(smoking~.,  
                        data = train)  
  
varImpPlot(rf_init)
```

rf_init

hemoglobin
height.cm.
Gtp
triglyceride
LDL
HDL
waist.cm.
Cholesterol
ALT
serum.creatinine
fasting.blood.sugar
weight.kg.
systolic
AST
age
relaxation
eyesight.right.
eyesight.left.
dental.caries
Urine.protein
hearing.right.
hearing.left.



rf_init

```
##
## Call:
## randomForest(formula = smoking ~ ., data = train)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 4
##
## OOB estimate of  error rate: 22.67%
## Confusion matrix:
##      0      1 class.error
## 0 29349  9968  0.2535290
## 1  5900 24783  0.1922889
```

Here we see that our initial model is built with a default of 500 trees and 4 variables in each tree.

We also see that based on the mean decrease in the gini impurity, the most important variables in this classification are: hemoglobin, height, Gtp and triglyceride. There is a pretty sharp decrease in importance once we get to LDL.

Mean Decrease in Gini Impurity

Gini Impurity is a measure commonly employed in decision tree algorithms, particularly in CART (Classification and Regression Tree). It quantifies the likelihood that a randomly chosen observation would be incorrectly classified if it were randomly labeled based on the distribution of categories in a set of data. The Gini Impurity ranges from 0 to 0.5, where 0 indicates perfect purity (all cases belong to a single category), and 0.5 implies that the classification is as good as a random guess. In other words, a lower Gini Impurity signifies a more homogeneous node in the decision tree, contributing to a more effective split during the tree-building process.

Finding the Optimal Random Forest Parameters

With these facts in mind, let's find the optimal parameters for our Random Forest model. We will be optimizing the number of 'weak' trees 'grown' in the forest (ntree) and the number of independent variables included in each tree (mtry).

We will build 16 Random Forest models. The goal is to exhaust every desired ntree and mtry value then choose the best model.

```
k <- 1
rf_list <- list()
ntree_param <- c(400,500,600, 700)
mtry_param <- c(3,4,5,6)

for(trees in ntree_param){
  for(vars in mtry_param){
    model <- randomForest(smoking~.,
                          data = train,
                          ntree = trees,
                          mtry = vars)

    rf_list[[k]] <- model
    k = k+1
  }
}
```

Choosing the Optimal Model


We will be using the Out-Of-Bag Error to find the best model. OOB Error also utilizes Bootstrap Aggregation. When we create those previously mentioned 'weak' trees using bootstrap samples, we can calculate an almost pseudo-out-of-sample error rate for each tree using the data that was NOT randomly selected to train the tree. How well will the tree predict the observations that it was not trained on... but that are still within the training data?

We will be looking at the OOB Error Rate for each of the 16 models and choosing the model with the lowest value.

```
rf_err_vec <- c()

for(i in 1:length(rf_list)){
  rf_err_vec[i] <- tail(rf_list[[i]]$err.rate[, 'OOB'], 1)
}

plot(rf_err_vec)
```



```
which.min(rf_err_vec)
```

```
## [1] 14
```

```
rf_list[[which.min(rf_err_vec)]]
```



```
##
## Call:
## randomForest(formula = smoking ~ ., data = train, ntree = trees,      mtry = vars
)
##              Type of random forest: classification
##              Number of trees: 700
## No. of variables tried at each split: 4
##
##              OOB estimate of  error rate: 22.53%
## Confusion matrix:
##           0      1 class.error
## 0 29413  9904   0.2519012
## 1  5868 24815   0.1912460
```

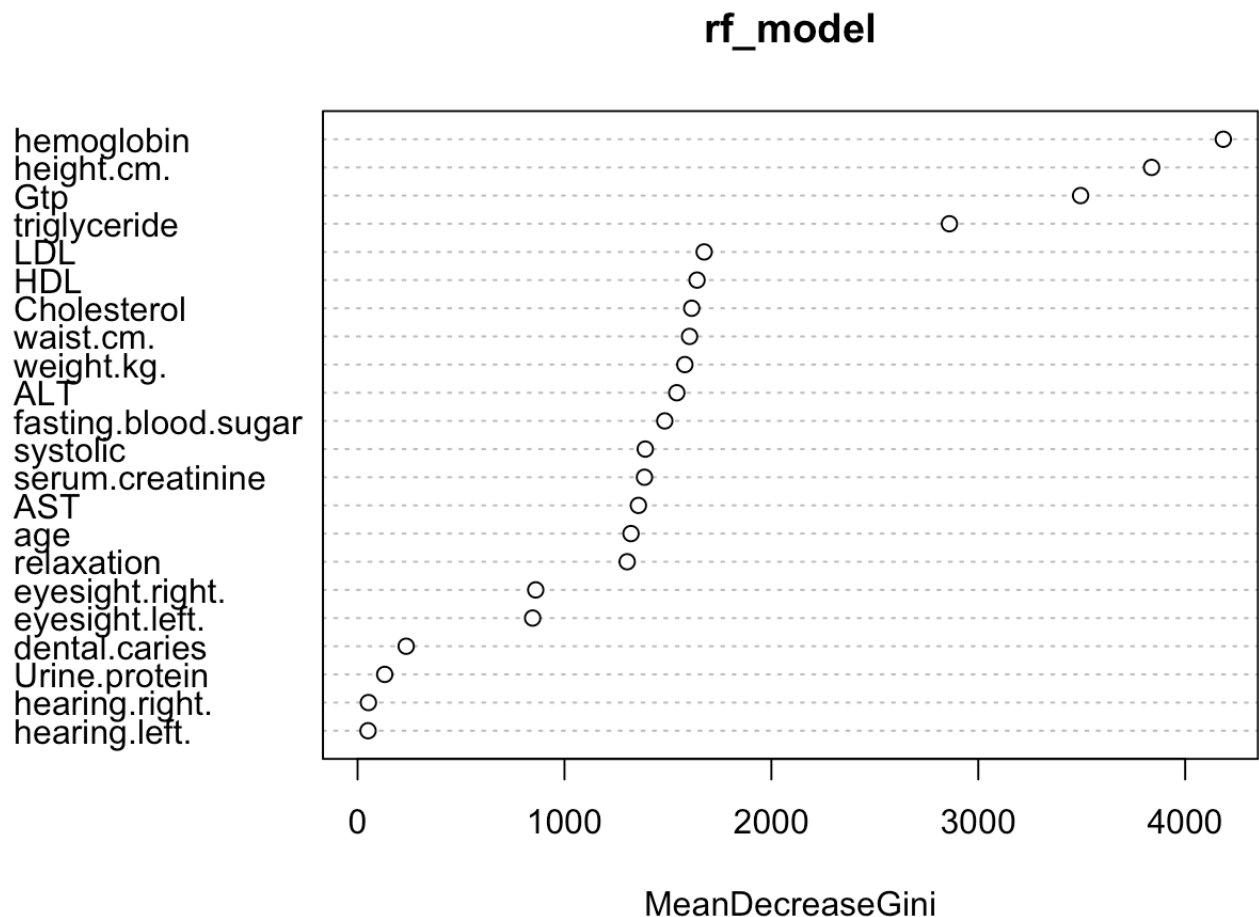
```
rf_model <- rf_list[[which.min(rf_err_vec)]]

rf_model
```

```
##
## Call:
## randomForest(formula = smoking ~ ., data = train, ntree = trees,      mtry = vars
)
##              Type of random forest: classification
##              Number of trees: 700
## No. of variables tried at each split: 4
##
##              OOB estimate of  error rate: 22.53%
## Confusion matrix:
##           0      1 class.error
## 0 29413  9904   0.2519012
## 1  5868 24815   0.1912460
```

Because Random Forest uses bootstrapping and random sampling, every time we run the algorithm, we get different results. The parameters for the optimal tree (with the lowest OOB Error Rate) will be shown above.

```
varImpPlot(rf_model)
```



The variable importance has not changed much from our initial random forest model.

Logistic Regression

Now that we have an idea of what our important variables are based on Gini Impurity, we can begin thinking about fitting a logistic regression model.


Logistic regression uses the sigmoid function (an S-shaped curve) and assigns probabilities to points on this curve. It is commonly used for classification problems. Logistic regression uses the MLE (Maximum Likelihood Estimate) method to estimate the parameters of an assumed probability distribution by maximizing the likelihood function.

To strengthen the power of this model, I'd like to include interaction terms. This is typically done via domain knowledge and is difficult to do via statistical tests. Using a stepwise logit regression is an option, but is computationally inefficient (it would take hours). Another option (not exactly perfect either) is to revert to visualizations. After iterating through many scatterplots, we find two possible interactions to include in our model.

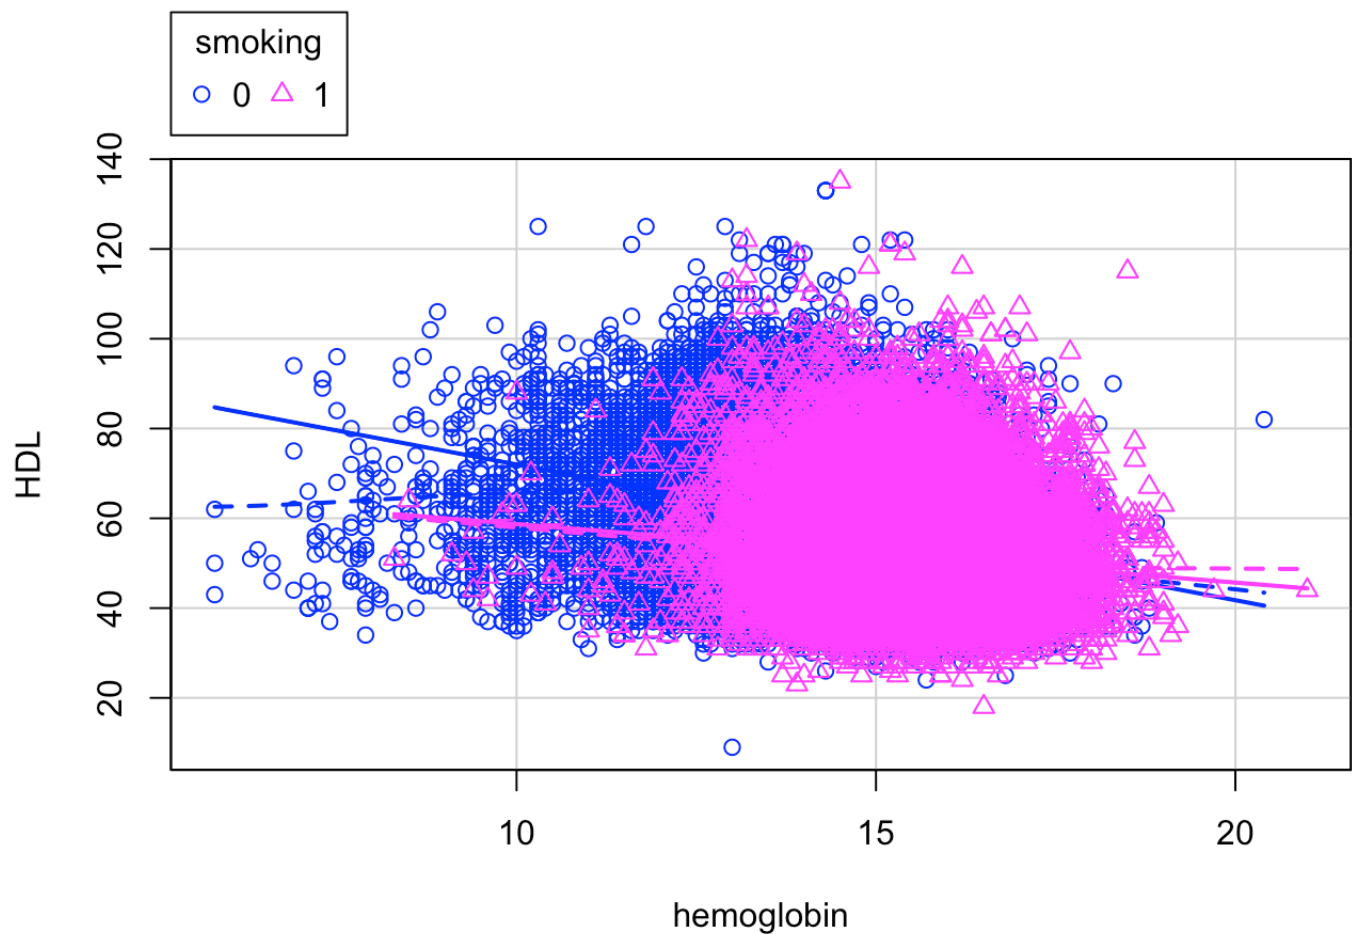
```
library(car)
```

```
## Loading required package: carData
```

```
scatterplot(formula = Cholesterol~hemoglobin | smoking , data=train)
```



```
scatterplot(formula = HDL~hemoglobin | smoking , data=train)
```



Let's fit our first logit model with the previously found 'important variables' and these two interaction terms.

```
logit.1 <- glm(formula = smoking~hemoglobin+height.cm.+Gtp+triglyceride+LDL+HDL+Cholesterol+Cholesterol:hemoglobin+HDL:hemoglobin,
               data = train,
               family = binomial)
summary(logit.1)
```

```
##
## Call:
## glm(formula = smoking ~ hemoglobin + height.cm. + Gtp + triglyceride +
##      LDL + HDL + Cholesterol + Cholesterol:hemoglobin + HDL:hemoglobin,
##      family = binomial, data = train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -5.3277692   0.9660932  -5.515 3.49e-08 ***
## hemoglobin      -0.4467334   0.0621290  -7.190 6.46e-13 ***
## height.cm.       0.0737653   0.0013853  53.250 < 2e-16 ***
## Gtp              0.0167457   0.0004262  39.290 < 2e-16 ***
## triglyceride     0.0071351   0.0002430  29.360 < 2e-16 ***
## LDL             -0.0008717   0.0007837  -1.112  0.266
## HDL             -0.1246563   0.0100324 -12.425 < 2e-16 ***
## Cholesterol     -0.0438327   0.0044826  -9.778 < 2e-16 ***
## hemoglobin:Cholesterol 0.0022996  0.0002904   7.920 2.38e-15 ***
## hemoglobin:HDL    0.0082156  0.0006633  12.386 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 95973  on 69999  degrees of freedom
## Residual deviance: 69557  on 69990  degrees of freedom
## AIC: 69577
##
## Number of Fisher Scoring iterations: 5
```

Each variable looks to be significant!

Logit 2

Now, just to be safe, let's run a logit model with only the main effects (no interaction terms).

```
logit.2 <- glm(formula = smoking~hemoglobin+height.cm.+Gtp+triglyceride+LDL+HDL+Chole
sterol,
               data = train,
               family = binomial)
summary(logit.2)
```

```
##
## Call:
## glm(formula = smoking ~ hemoglobin + height.cm. + Gtp + triglyceride +
##      LDL + HDL + Cholesterol, family = binomial, data = train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.894e+01  2.393e-01 -79.142  <2e-16 ***
## hemoglobin   4.411e-01  9.264e-03  47.611  <2e-16 ***
## height.cm.   7.541e-02  1.385e-03  54.448  <2e-16 ***
## Gtp          1.704e-02  4.261e-04  39.993  <2e-16 ***
## triglyceride 7.090e-03  2.444e-04  29.010  <2e-16 ***
## LDL          -9.219e-04  7.939e-04  -1.161    0.246
## HDL          -1.928e-03  1.174e-03  -1.642    0.101
## Cholesterol -9.169e-03  8.685e-04 -10.558  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 95973  on 69999  degrees of freedom
## Residual deviance: 69811  on 69992  degrees of freedom
## AIC: 69827
##
## Number of Fisher Scoring iterations: 5
```

Once again, each variable is significant, but our AIC has increased.

Model Evaluation

Accuracy

A model's accuracy measures it's ability to detect correct positive and negative cases among all actual positive and negative cases.

Random Forest

The in-sample and out-of-sample accuracy's for our random forest model are fairly similar. However, our out-of-sample is a little higher. A great sign!

```
pred_in_rf <- predict(rf_model, type = 'response')
pred_out_rf <- predict(rf_model, newdata = valid, type = 'response')

# in-sample accuracy
sum(diag(rf_model$confusion[,c(1,2)]))/sum(rf_model$confusion[,c(1,2)])
```

```
## [1] 0.7746857
```

```
# out-of-sample accuracy
rf_out_confusion <- table(valid$smoking, pred_out_rf)
sum(diag(rf_out_confusion))/sum(rf_out_confusion)
```

```
## [1] 0.7765977
```

Logit

With both of our logistic regression models, our out-of-sample accuracy outperforms our in-sample accuracy. Logit.1 (Logistic Regression Model with Interaction Terms) is performing just a little better than Logit.2 (Logistic Regression Model with only Main Effects).

```
pred_in_logit1 <- predict(logit.1, type = 'response')
pred_out_logit1 <- predict(logit.1, newdata = valid, type = 'response')
logit1_in_confusion <- table(train$smoking, ifelse(pred_in_logit1>0.5,1,0))
logit1_out_confusion <- table(valid$smoking, ifelse(pred_out_logit1>0.5,1,0))
# in-sample logit.1
sum(diag(logit1_in_confusion))/sum(logit1_in_confusion)
```

```
## [1] 0.7403286
```

```
# out-of-sample logit.1
sum(diag(logit1_out_confusion))/sum(logit1_out_confusion)
```

```
## [1] 0.7416308
```

```
pred_in_logit2 <- predict(logit.2, type = 'response')
pred_out_logit2 <- predict(logit.2, newdata = valid, type = 'response')
logit2_in_confusion <- table(train$smoking, ifelse(pred_in_logit2>0.5,1,0))
logit2_out_confusion <- table(valid$smoking, ifelse(pred_out_logit2>0.5,1,0))
# in-sample logit.2
sum(diag(logit2_in_confusion))/sum(logit2_in_confusion)
```

```
## [1] 0.7393714
```

```
# out-of-sample logit.2
sum(diag(logit2_out_confusion))/sum(logit2_out_confusion)
```

```
## [1] 0.7408241
```

Sensitivity

A model's sensitivity measures its ability to detect positive cases among all actual positive cases.

Our Random Forest model outperforms the Logistic Regression models in the sensitivity measure.

```
# out-of-sample RF
sum(ifelse(valid$smoking==1 & pred_out_rf==1,1,0))/sum(valid$smoking==1)
```

```
## [1] 0.8109571
```

```
# out-of-sample logit.1
sum(ifelse(valid$smoking==1 & pred_out_logit1>0.5,1,0))/sum(valid$smoking==1)
```

```
## [1] 0.7334873
```

```
# out-of-sample logit.2
sum(ifelse(valid$smoking==1 & pred_out_logit2>0.5,1,0))/sum(valid$smoking==1)
```

```
## [1] 0.7260457
```

Specificity

A model's specificity measures its ability to detect negative cases among all actual negative cases.

In specificity, our logit.2 (Logistic Regression with only Main Effects) model outperforms all other models.

```
# out-of-sample RF
sum(ifelse(valid$smoking==0 & pred_out_rf==0,1,0))/sum(valid$smoking==0)
```

```
## [1] 0.7499702
```

```
# out-of-sample logit.1
sum(ifelse(valid$smoking==0 & pred_out_logit1<=0.5,1,0))/sum(valid$smoking==0)
```

```
## [1] 0.7479418
```

```
# out-of-sample logit.2
sum(ifelse(valid$smoking==0 & pred_out_logit2<=0.5,1,0))/sum(valid$smoking==0)
```

```
## [1] 0.752277
```

ROC Curves and AUC


Receiver Operating Characteristic (ROC) curves plot the performance of binary classification models. They illustrate the relationship between the True Positive Rate (Sensitivity) and the False Positive Rate (1 - Specificity) at varying threshold values for classifying the positive class. These thresholds are the points where the model transitions between designating a case as True or False for the dependent variable, often based on predicted probabilities.

The Area Under the Curve (AUC) represents the area under the ROC curve. A higher AUC generally indicates a stronger model, but it's important to note that a perfect model has an AUC of 1.0, and a random or ineffective model has an AUC of 0.5. Therefore, higher AUC values signify better discriminative ability.

```
prob_out_rf_prob <- predict(rf_model, newdata = valid, type = 'prob')[,2]

# Random Forest
roc.curve(response = valid$smoking, predicted = prob_out_rf_prob)
```


ROC curve



```
## Area under the curve (AUC): 0.858
```

```
# Logit 1
roc.curve(response = valid$smoking, predicted = pred_out_logit1)
```


ROC curve



```
## Area under the curve (AUC): 0.825
```

```
# Logit 2  
roc.curve(response = valid$smoking, predicted = pred_out_logit2)
```

ROC curve



```
## Area under the curve (AUC): 0.825
```

Our Random Forest model has the largest Area Under the Curve, indicating that it has the best discriminative power out of the three models.

While our Logistic Regression models demonstrated competitive performance, the Random Forest approach showcased its effectiveness in leveraging diverse bio-signals for accurate identification of smokers and non-smokers. This underscores the potential of machine learning, particularly Random Forest, in elucidating the nuanced relationship between bio-signals and smoking habits for enhanced predictive modeling.